

EIE 3112

Relational Models and ER Diagrams

Dr Pauli Lai

T. Connolly and C. Begg, “*Database Systems: A Practical Approach to Design, Implementation, and Management*,” 6th Edition, Chapter 4, Pearson, 2015. (5th Edition is also fine)

Intended Learning Outcomes (ILO)

Subject ILO to achieve:

A: Professional/academic knowledge and skills

1. Database design, development, and programming

Topics:

- ◆ Terminology of relational model.
- ◆ How tables are used to represent data.
- ◆ How to identify CK, PK, and FKs.
- ◆ Meaning of entity integrity and referential integrity
- ◆ ER Diagrams

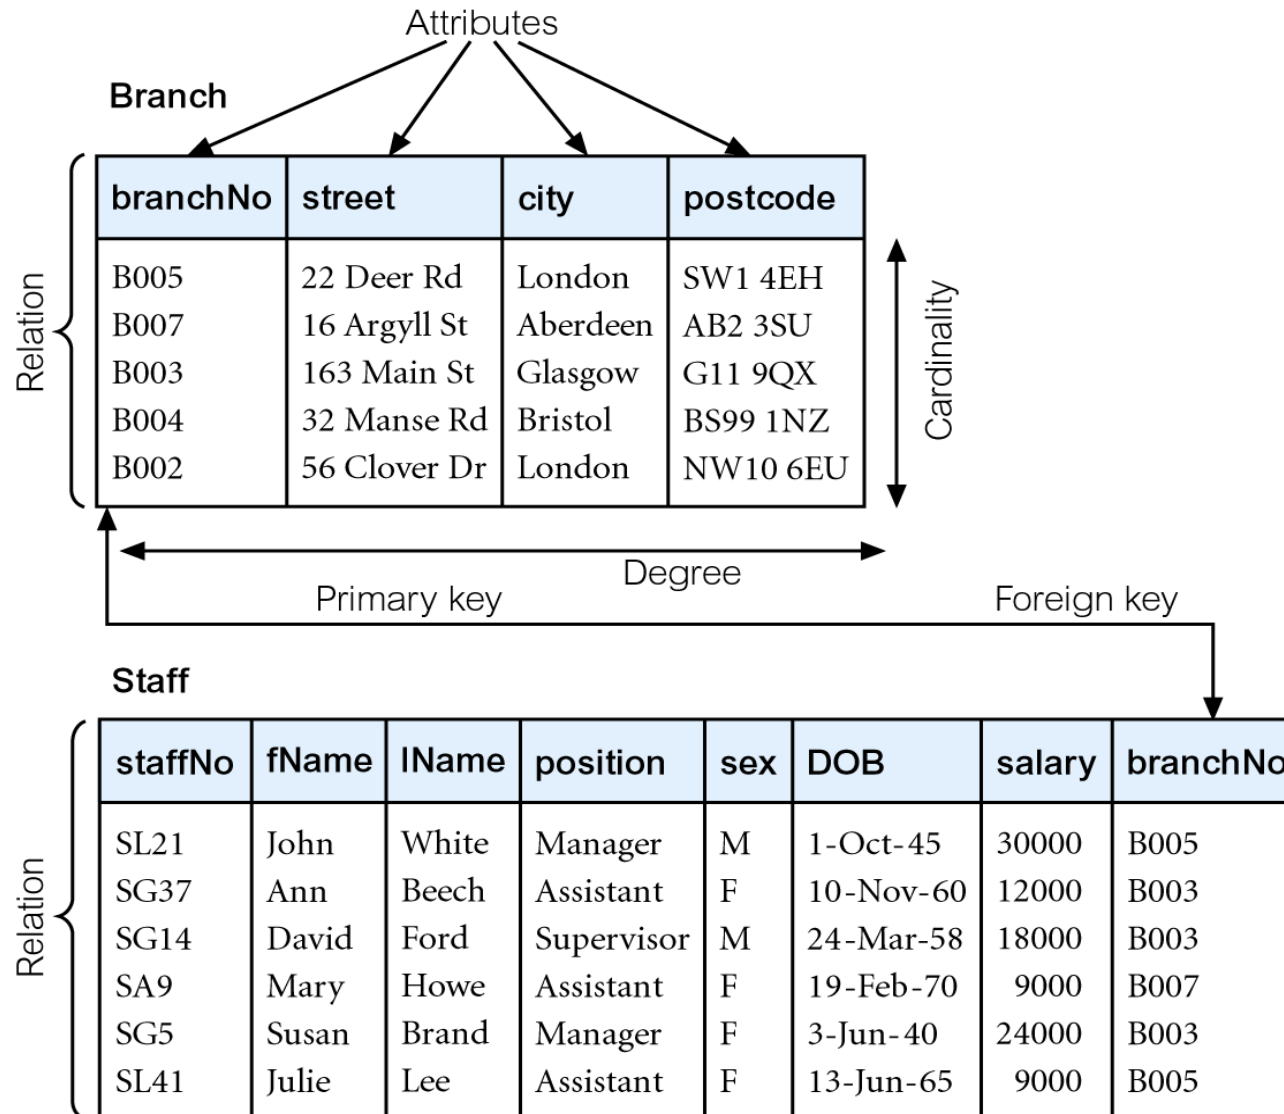
Relational Model Terminology

- ◆ A relation is a table with columns and rows.
 - ◆ Only applies to logical structure of the database, not the physical structure.
- ◆ Attribute is a named column of a relation.
- ◆ Domain is the set of allowable values for one or more attributes.

Relational Model Terminology

- ◆ Tuple is a row of a relation.
- ◆ Degree is the number of attributes in a relation.
- ◆ Cardinality is the number of tuples in a relation.
- ◆ Relational Database is a collection of relations with distinct relation names.

Branch and Staff Relations



Example of Attribute Domains

Attribute	Domain Name	Meaning	Domain Definition
branchNo	BranchNumbers	The set of all possible branch numbers	character: size 4, range B001–B999
street	StreetNames	The set of all street names in Britain	character: size 25
city	CityNames	The set of all city names in Britain	character: size 15
postcode	Postcodes	The set of all postcodes in Britain	character: size 8
sex	Sex	The sex of a person	character: size 1, value M or F
DOB	DatesOfBirth	Possible values of staff birth dates	date, range from 1-Jan-20, format dd-mmm-yy
salary	Salaries	Possible values of staff salaries	monetary: 7 digits, range 6000.00–40000.00

Alternative Terminology

Formal terms	Alternative 1	Alternative 2
Relation	Table	File
Tuple	Row	Record
Attribute	Column	Field

Class Activity: Q&A

- When I call your name, please turn on your mic to answer the question.
- Example questions related to the relational model terminology:
 - How many tuples are there in relation Branch?
 - What is the Cardinality of table Staff?
 - How many attributes are there in file Staff?
 - What is the Degree of relation Branch?
 - How many fields are there in table Staff?
 - What is the domain for staffNo in relation Staff?
 - What is the domain for position in relation Staff?
 - Can you suggest one more column for Staff?
 - Can you suggest one more row for Branch?

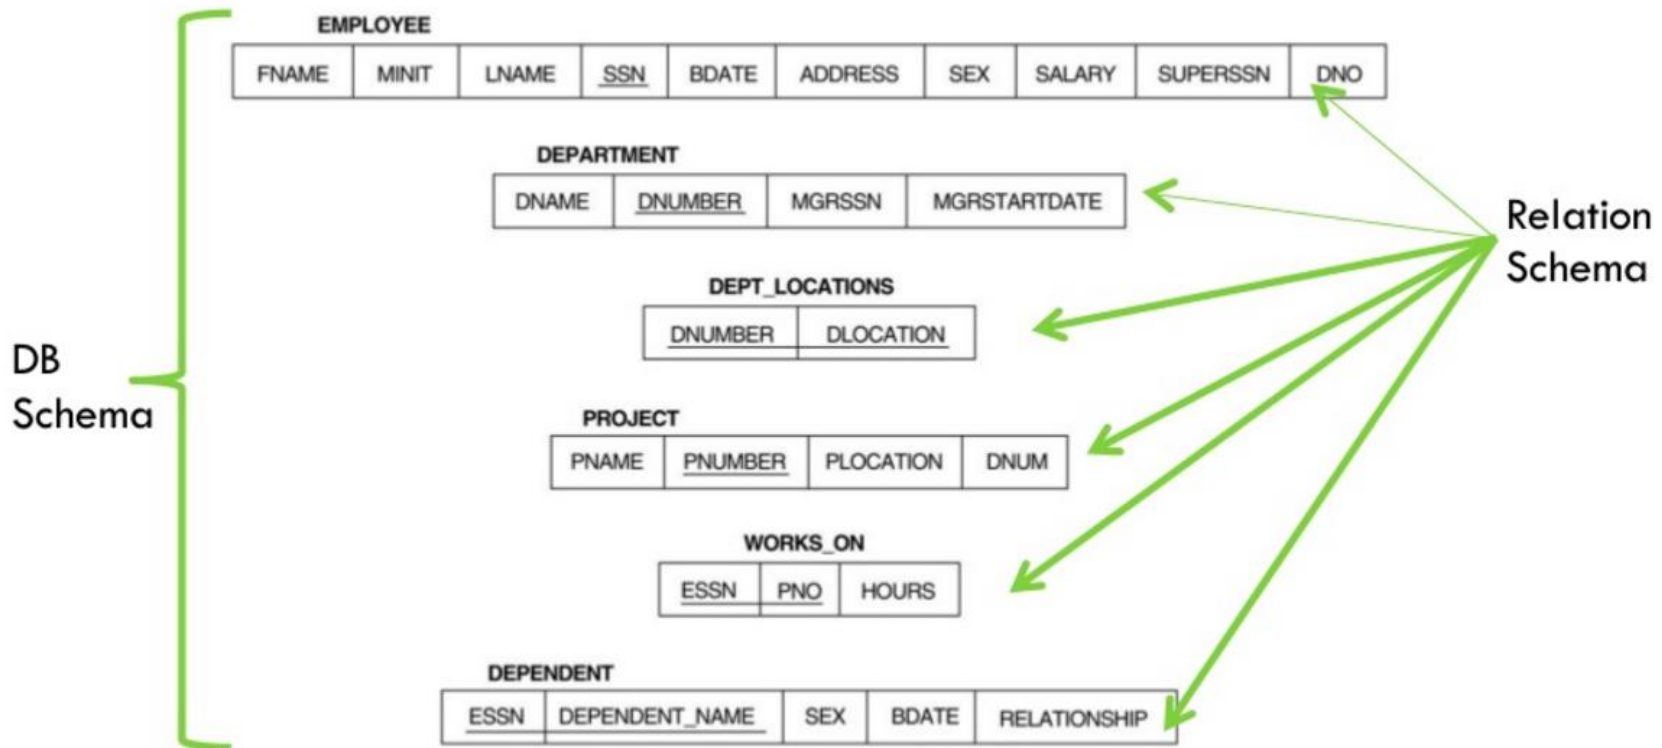
Database Relations

- ◆ Relation schema
 - ◆ Named relation defined by a set of attribute and domain name pairs.
- ◆ Relational database schema
 - ◆ Set of relation schemas, each with a distinct name.

Relational Database schema

26

- A sets of relation schema.
- The example shows the database schema for COMPANY = {EMPLOYEE, DEPARTMENT, DEPENDENT, PROJECT}



Properties of Relations

- ◆ Relation name is distinct from all other relation names in relational DB schema.
- ◆ Each cell of relation contains exactly one atomic (single) value.
- ◆ Each attribute has a distinct name.
- ◆ Values of an attribute are all from the same domain.

Properties of Relations

- ◆ Each tuple is distinct; there are no duplicate tuples.
- ◆ Order of attributes has no significance.
- ◆ Order of tuples has no significance, theoretically.

Relational Keys

◆ Superkey

- ◆ An attribute, or set of attributes, that uniquely identifies a tuple within a relation.
- ◆ A superkey may contain additional attributes that are not necessary for unique identification.

Superkey?

Superkey?

staffNo	fName	lName	position	sex	DOB	salary	branchNo
SL21	John	White	Manager	M	1-Oct-45	30000	B005
SG37	Ann	Beech	Assistant	F	10-Nov-60	12000	B003
SG14	David	Ford	Supervisor	M	24-Mar-58	18000	B003
SA9	Mary	Howe	Assistant	F	19-Feb-70	9000	B007
SG5	Susan	Brand	Manager	F	3-Jun-40	24000	B003
SL41	Julie	Lee	Assistant	F	13-Jun-65	9000	B005

Superkey?

Relational Keys

◆ Candidate Key

- ◆ Superkey (K) such that no proper subset is a superkey within the relation (minimum superkey)
- ◆ In each tuple of relation R , values of candidate K uniquely identify that tuple (uniqueness).
- ◆ No proper subset of K has the uniqueness property (irreducibility).
- ◆ Identifying a candidate key requires that we know the “real-world” meaning of the attributes

Candidate Key? Candidate Key?

staffNo	fName	lName	position	sex	DOB	salary	branchNo
SL21	John	White	Manager	M	1-Oct-45	30000	B005
SG37	Ann	Beech	Assistant	F	10-Nov-60	12000	B003
SG14	David	Ford	Supervisor	M	24-Mar-58	18000	B003
SA9	Mary	Howe	Assistant	F	19-Feb-70	9000	B007
SG5	Susan	Brand	Manager	F	3-Jun-40	24000	B003
SL41	Julie	Lee	Assistant	F	13-Jun-65	9000	B005

Candidate Key?

Relational Keys

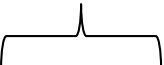
- **Primary Key**

- Candidate key selected to identify tuples uniquely within a relation.

- **Alternate Keys**

- Candidate keys that are not selected to be primary key.

Primary Key

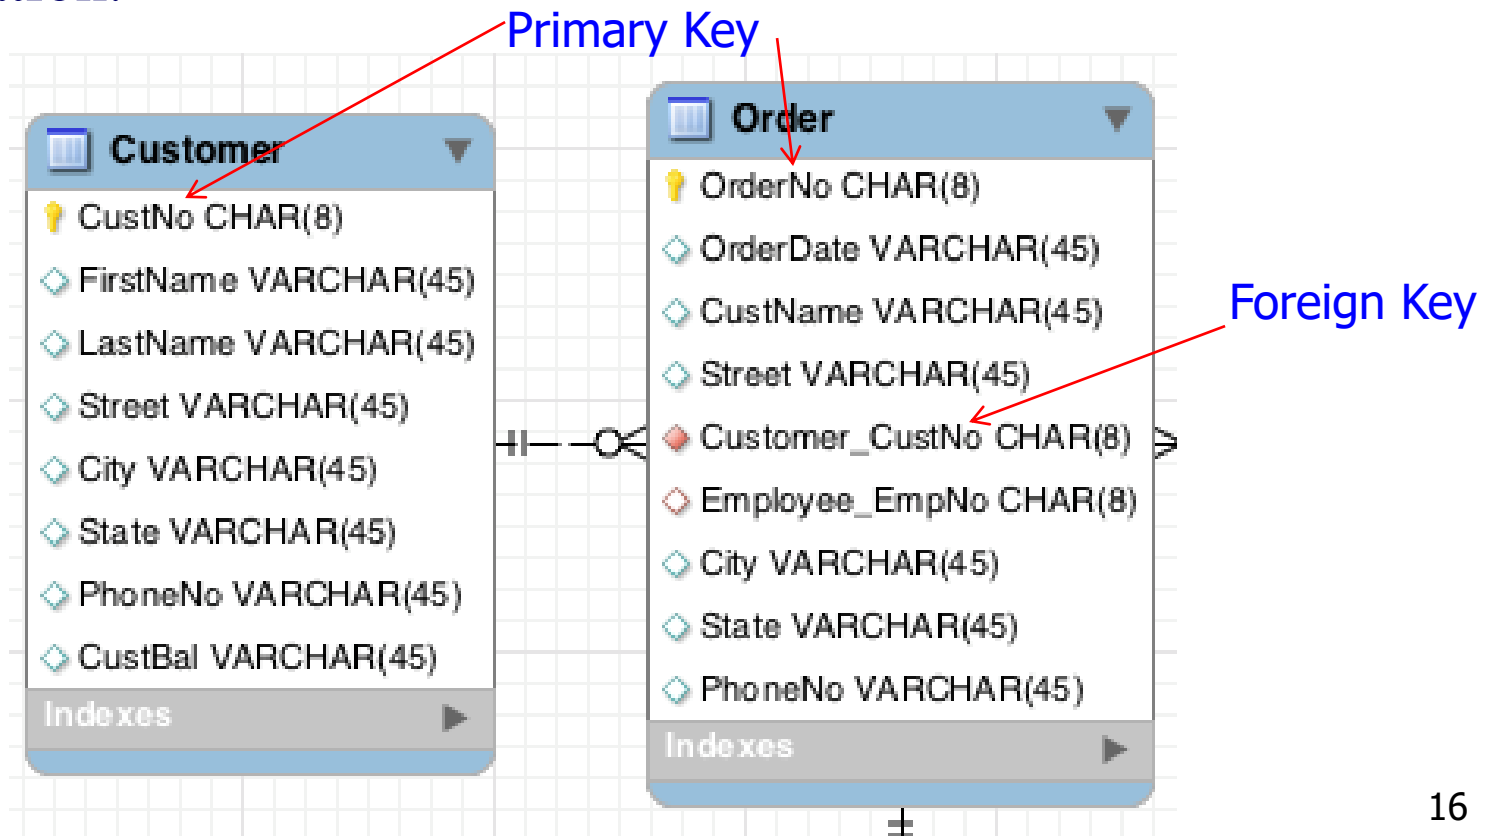


staffNo	fName	lName	position	sex	DOB	salary	branchNo
SL21	John	White	Manager	M	1-Oct-45	30000	B005
SG37	Ann	Beech	Assistant	F	10-Nov-60	12000	B003
SG14	David	Ford	Supervisor	M	24-Mar-58	18000	B003
SA9	Mary	Howe	Assistant	F	19-Feb-70	9000	B007
SG5	Susan	Brand	Manager	F	3-Jun-40	24000	B003
SL41	Julie	Lee	Assistant	F	13-Jun-65	9000	B005

Relational Keys

- **Foreign Key**

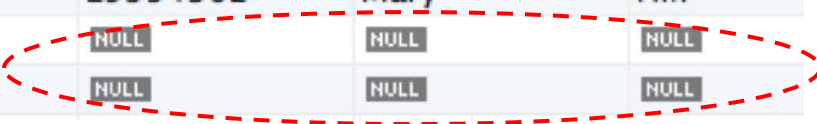
- Attribute, or set of attributes, within one relation that matches candidate key of some (possibly same) relation.



Integrity Constraints

- **Null**
 - Represents value for an attribute that is currently unknown or not applicable for tuple.
 - Deals with incomplete or exceptional data.
 - Represents the absence of a value and is not the same as zero or spaces, which are values.

OrderNo	OrderDate	EmpNo	FirstName	LastName
O1116324	1/23/2007	E8544399	Joe	Jenkins
O1231231	1/23/2007	E9954302	Mary	Hill
O1615141	1/23/2007	E8544399	Joe	Jenkins
O3252629	1/23/2007	E9954302	Mary	Hill
O7847172	1/23/2007	NULL	NULL	NULL
O8979495	1/23/2007	NULL	NULL	NULL



Integrity Constraints

- **Entity Integrity**

- Each relation has column(s) with unique values
- No two rows have the same value
- No attribute of a primary key can be null
- Can be achieved by having a primary key

- **Referential Integrity**

- Column values in one table must match column values in a related table.

Views

- **Base Relation**

- Named relation corresponding to an entity in conceptual schema, whose tuples are physically stored in database.

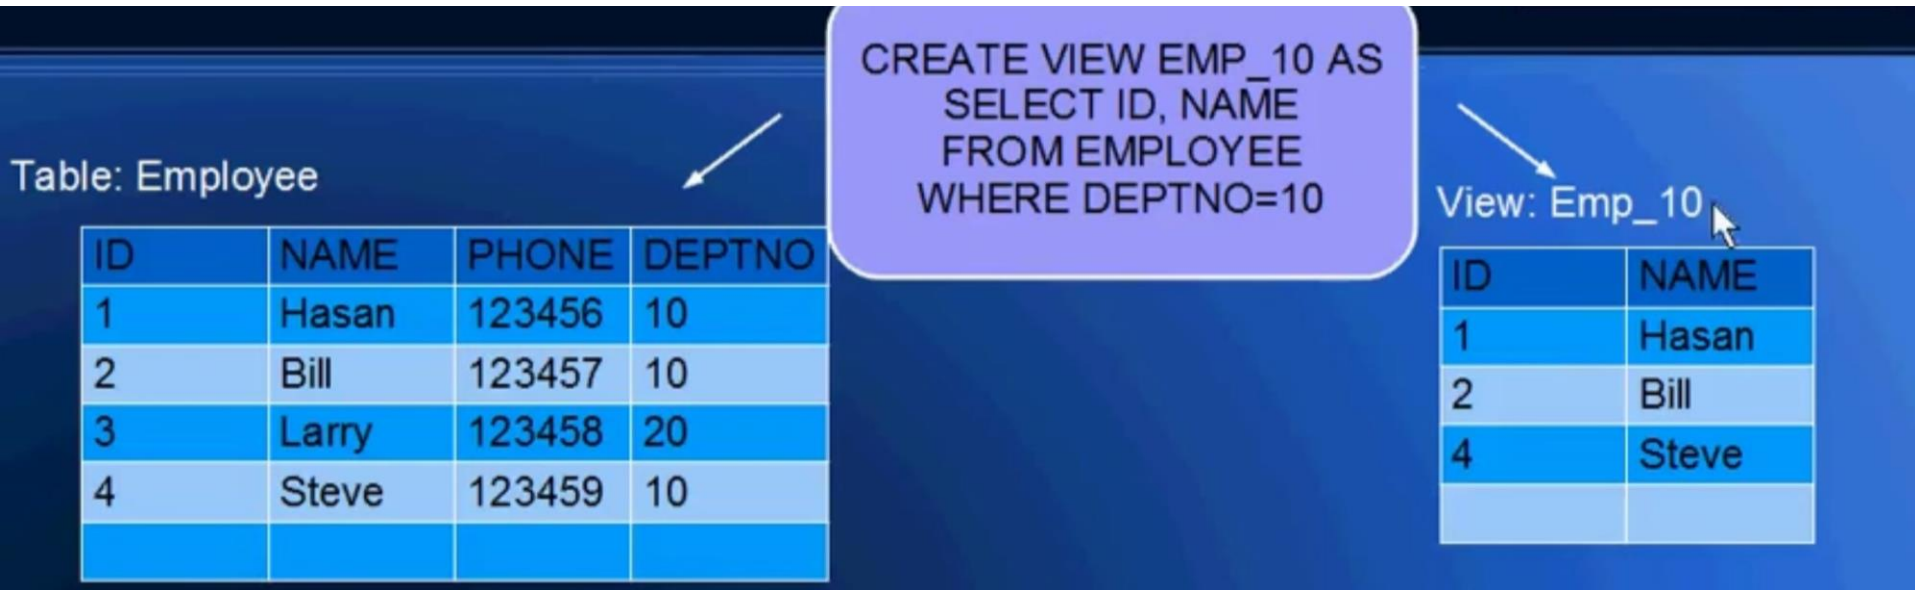
- **View**

- Dynamic result of one or more relational operations operating on base relations to produce another relation.

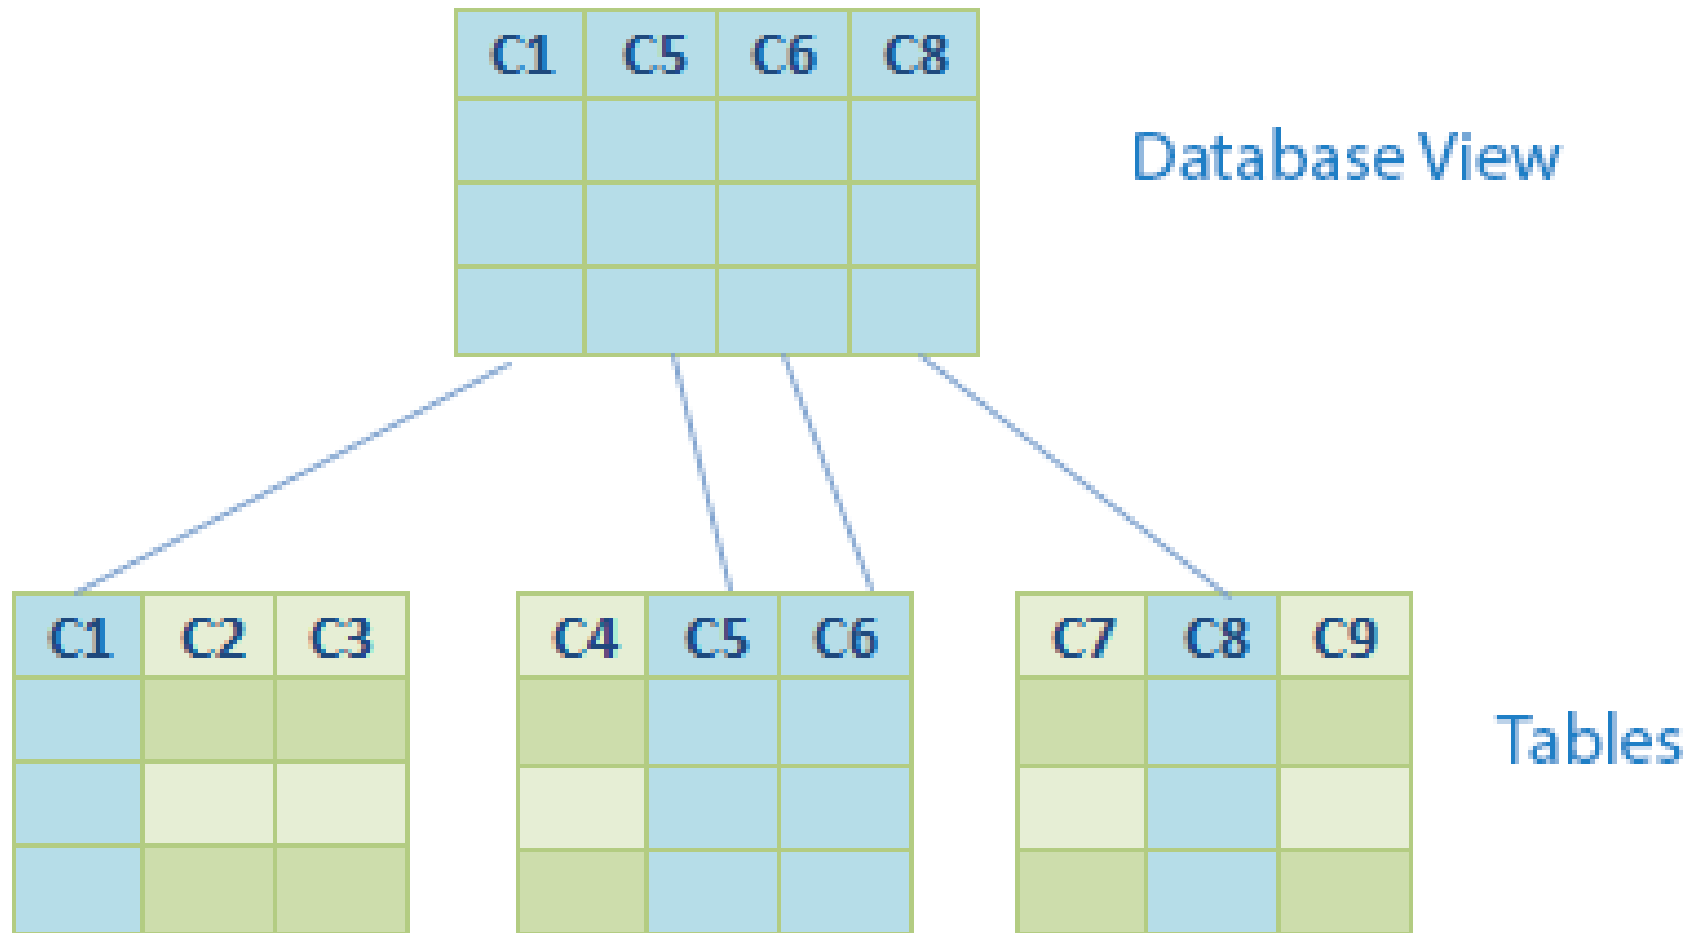
Views

Allows each user to have his or her own view of the database.

A view is essentially some subset of the database.



Views



View Types Based on Purpose

Synonym of a table

Subset of a table in terms
of rows or columns or both

Summarization of table's data

Join of multiple tables

Views - Benefits

- **Reduce complexity**
- **Provide a level of security**
- **Provide a mechanism to customize the appearance of the database**
- **Present a consistent, unchanging picture of the structure of the database, even if the underlying database is changed**

Updating Views

- All updates to a base relation should be immediately reflected in all views that reference that base relation.
- If view is updated, underlying base relation should reflect change.

Updating Views

- There are restrictions on types of modifications that can be made through views:
 - Updates are allowed if query involves a single base relation and contains a candidate key of base relation.
 - Updates are not allowed involving multiple base relations.
 - Updates are not allowed involving aggregation or grouping operations.
- You will learn more about Views in SQL lecture

Outline of ERD

- Conceptual data modeling
 - ERD: entity-relationship diagram
- Notations
 - Entities, relationships, cardinalities
- Relationships
 - One-to-one, one-to-many, many-to-many
 - Identifying Relationship

ER Diagram Terminologies

- Entities, attributes, identifiers, and relationships
- Entity:
 - Thing or object, whether real or imagined, about which information needs to be known/tracked
 - Examples:
 - Customer
 - Order
 - Student
 - Book
 - Employee

Entity

- Entity class/type:
 - A collection that share common properties or characteristics
 - UPPER CASE (noun)
- Entity instance:
 - Representation of a particular entity
- Many instances in an entity class

CUSTOMER entity class
contains:

CustID
FirstName
LastName
phoneNo

Two instances:

1	Peter	Lam	2134
2	Jenny	Chan	1234

Attributes

- Attributes:
 - Properties
 - describe the entity's characteristics

The diagram shows a table with four columns: CustID, FirstName, LastName, and phoneNo. The first two rows are highlighted in yellow. Annotations include: a blue arrow pointing to the first row labeled 'instance'; a pink oval around the 'LastName' header with a blue arrow pointing to it labeled 'attribute'; a pink oval around the '1234' value with a blue arrow pointing to it labeled 'Value of the attribute "phoneNo"'; and a blue oval around the entire first row.

CustID	FirstName	LastName	phoneNo
1	Peter	Lam	2134
2	Jenny	Chan	1234

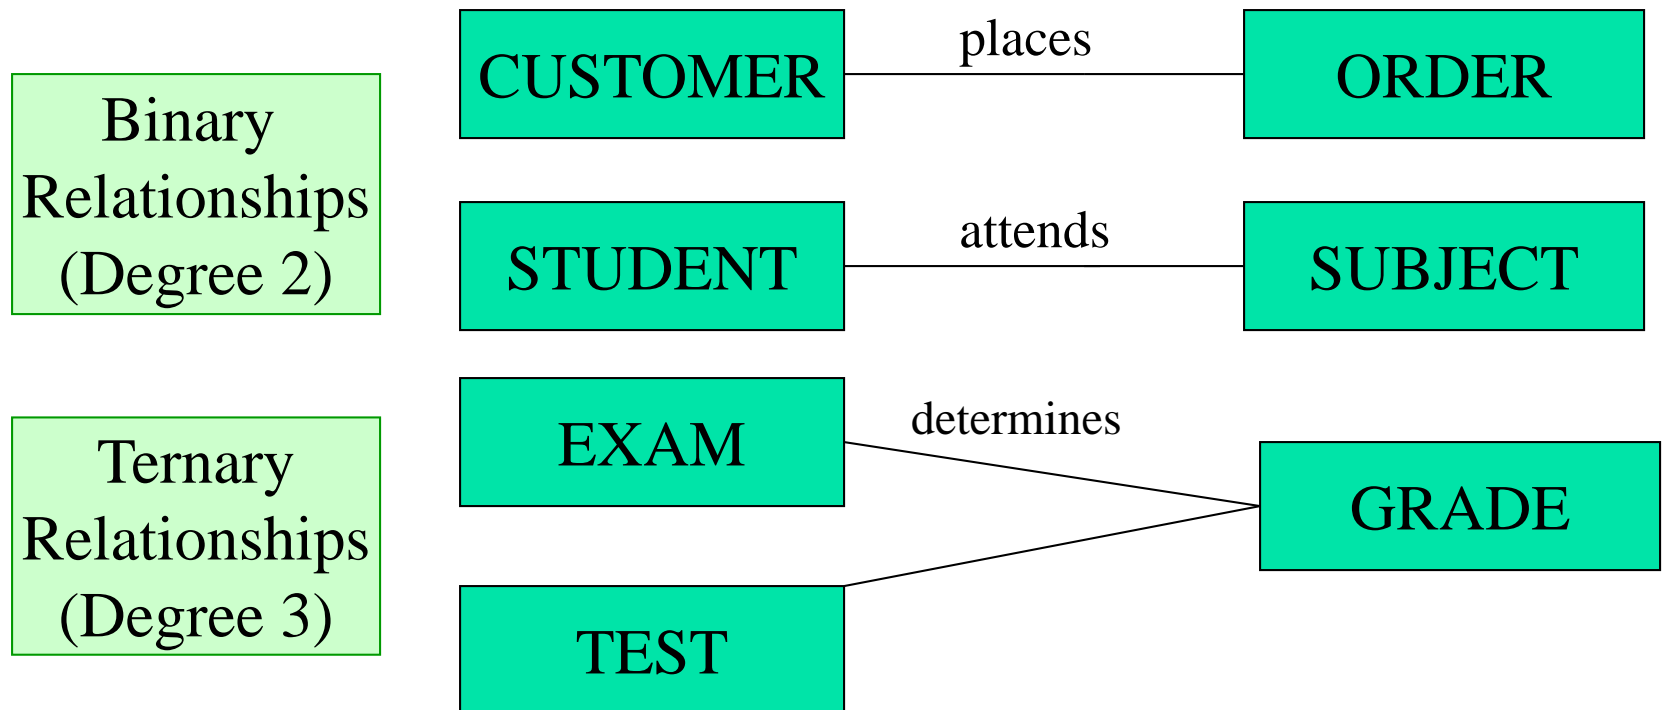
Identifiers

- Primary key
- Support entity identification
- One or more of the entity's attribute that uniquely identify a row in the table
 - CustID: uniquely identify a row
 - Not likely identified by LastName
 - Composite identifiers: identifier with 2 or more attributes

CustID	FirstName	LastName	phoneNo
1	Peter	Lam	2134
2	Jenny	Chan	1234

Relationships


- Describes how one or more entities are related to each other
- Entity: noun; Relationship: verb




Relationships

- Describes how one or more entities are related to each other
- Entity: noun; Relationship: verb

Unary
Relationships
(Degree 1)

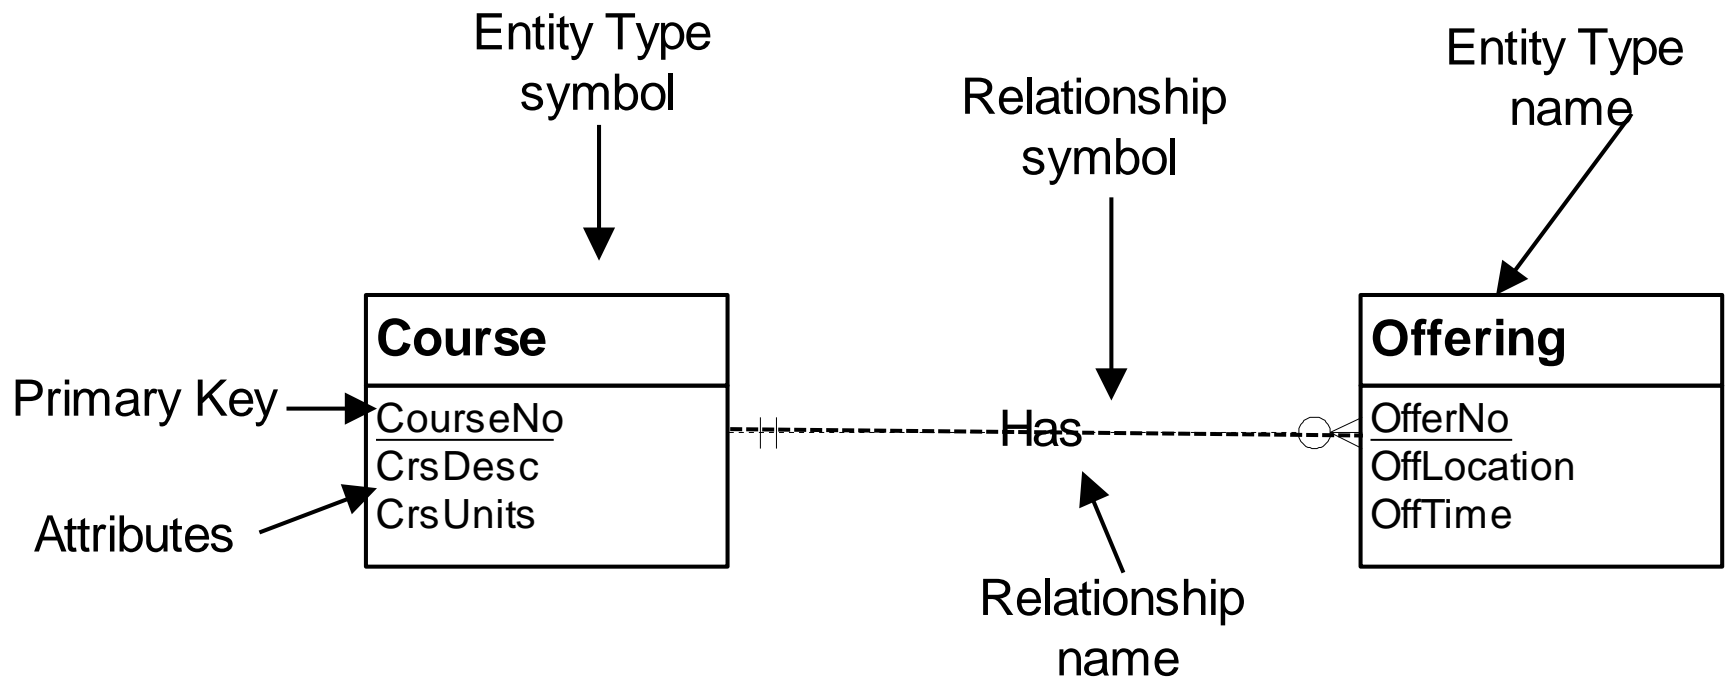


One-to-One



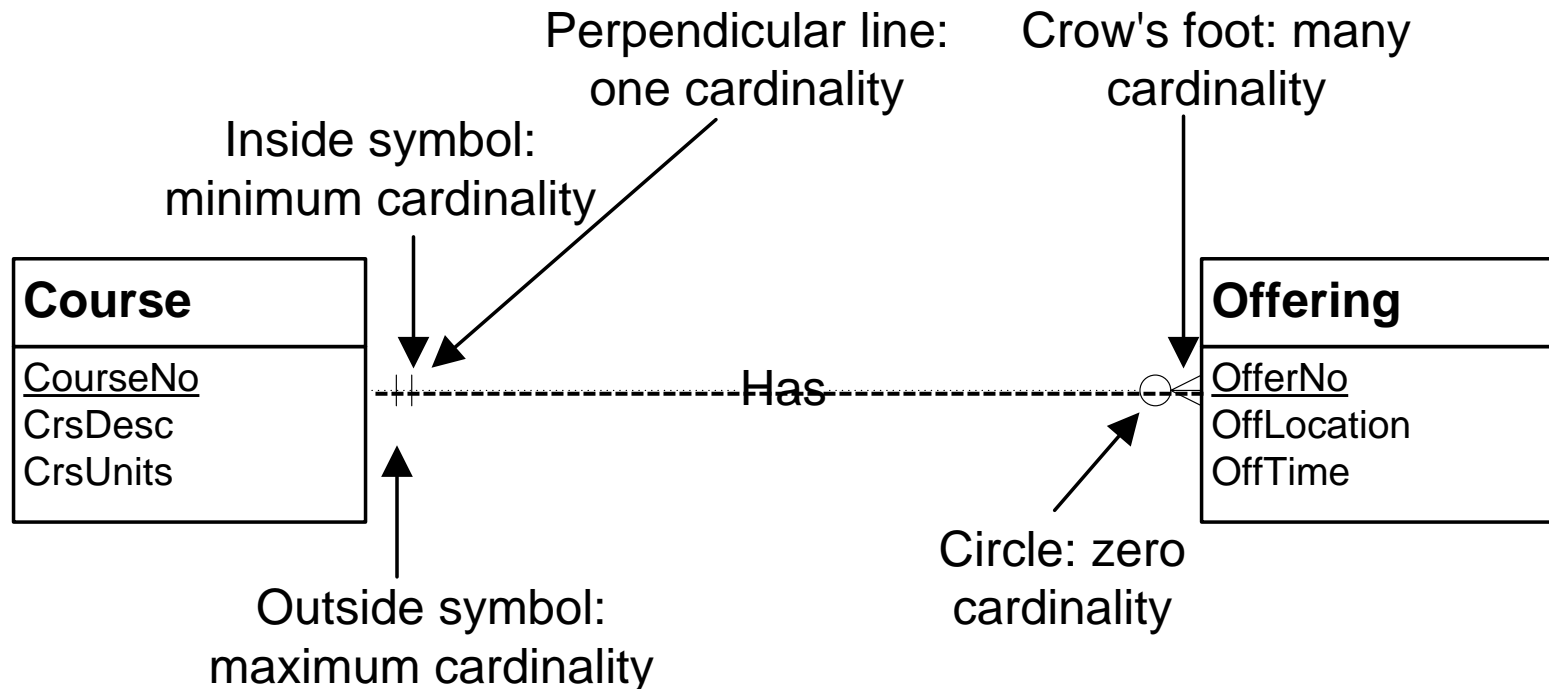
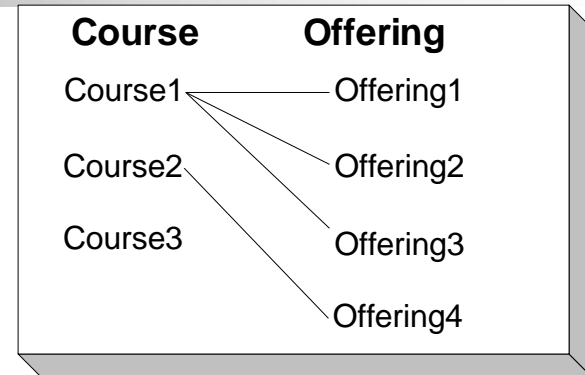
One-to-Many

ER Diagrams



Cardinalities

- Contains the number of objects that participate in a relationship



Cardinalities

- Minimum cardinality:
 - Indicates whether participation in the relationship is
 - Mandatory (1)
 - Optional (0)
- Maximum cardinality:
 - The maximum number of entities that can occur on one side of the relationship
 - One-to-one
 - One-to-many
 - Many-to-many

Cardinalities

Classification

Cardinality Restrictions

Mandatory

Minimum cardinality ≥ 1

Optional

Minimum cardinality = 0

Functional or single-valued

Maximum cardinality = 1

1-M

Maximum cardinality = 1 in one direction and
maximum cardinality > 1 in the other direction

M-N

Maximum cardinality is > 1 in both directions

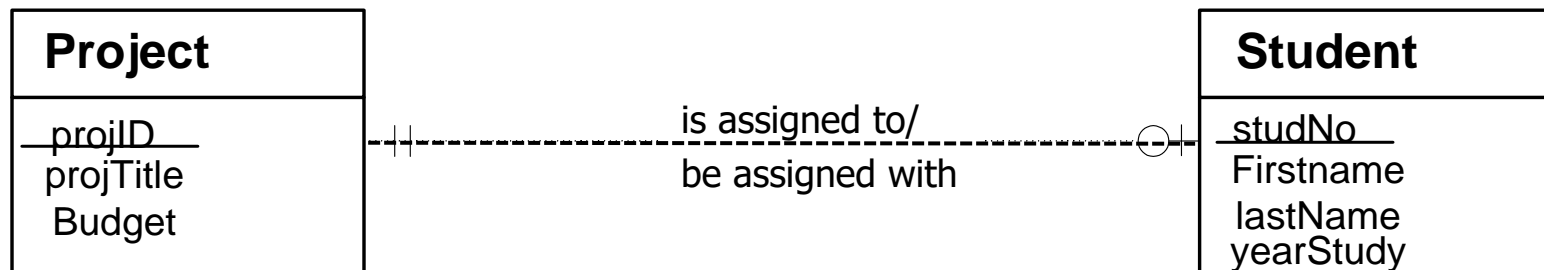
1-1

Maximum cardinality = 1 in both directions

Example 1

- **Project:**
 - Project title (**ProjTitle**), Project ID (**ProjID**) and Budget (**Budget**)
- **Student:**
 - Year of study (**YearStudy**), Student Number (**StudNo**), First Name (**FirstName**) and Last Name (**LastName**)
- **Relationship:**
 - A project might be assigned to at most one student
 - A student must be assigned one project

Example 1: answer

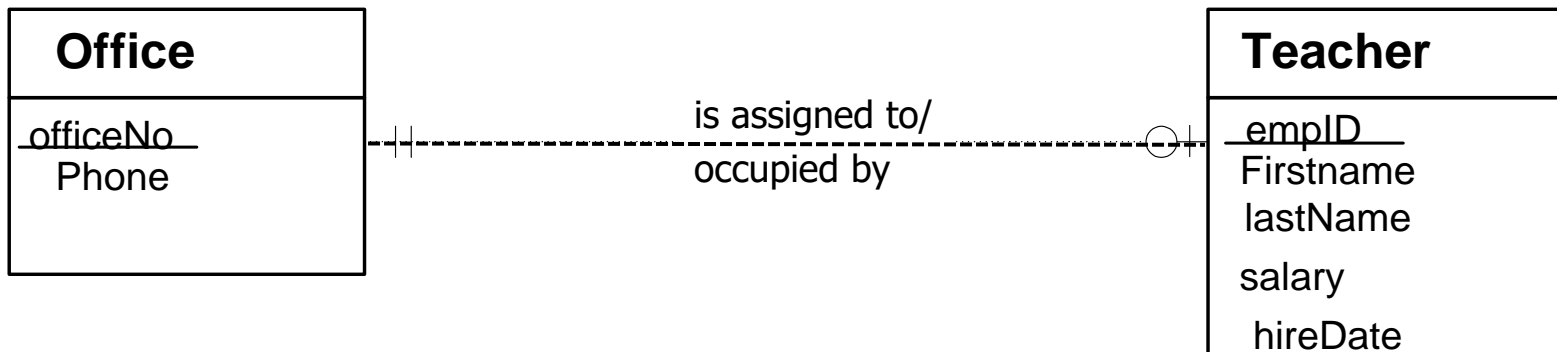


- A project might be assigned to at most one student
- A student must be assigned one project

Example 2

- Teacher:
 - Employee ID (ID), salary, firstname, lastname, hireDate
- Office:
 - officeNo, phone
- Each teacher is assigned to one office and an office is occupied by at most one teacher

Example 2: Answer

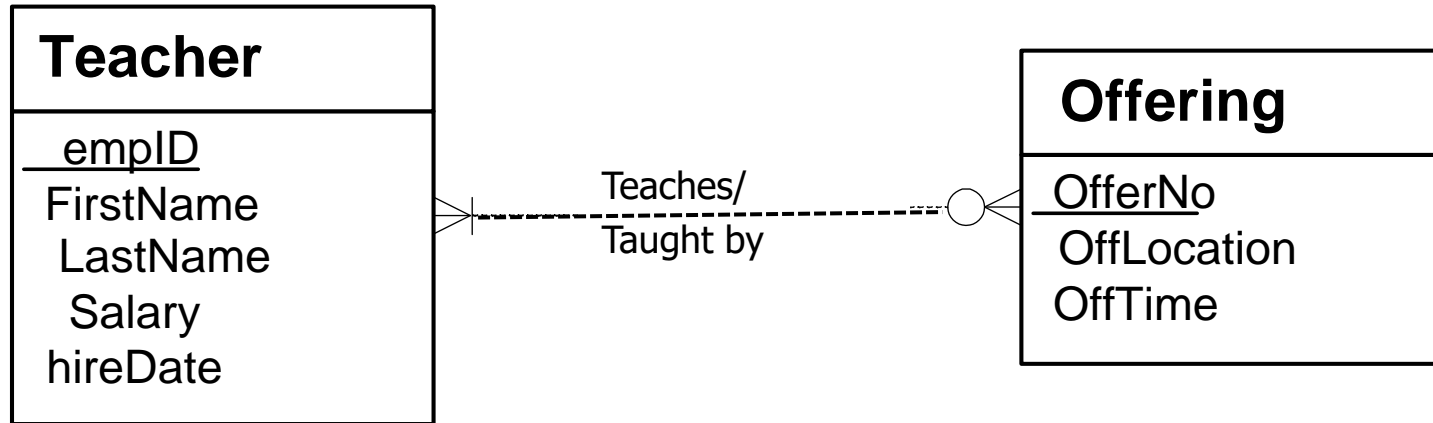


- Each teacher is assigned to one office and an office is occupied by at most one teacher

Example 3

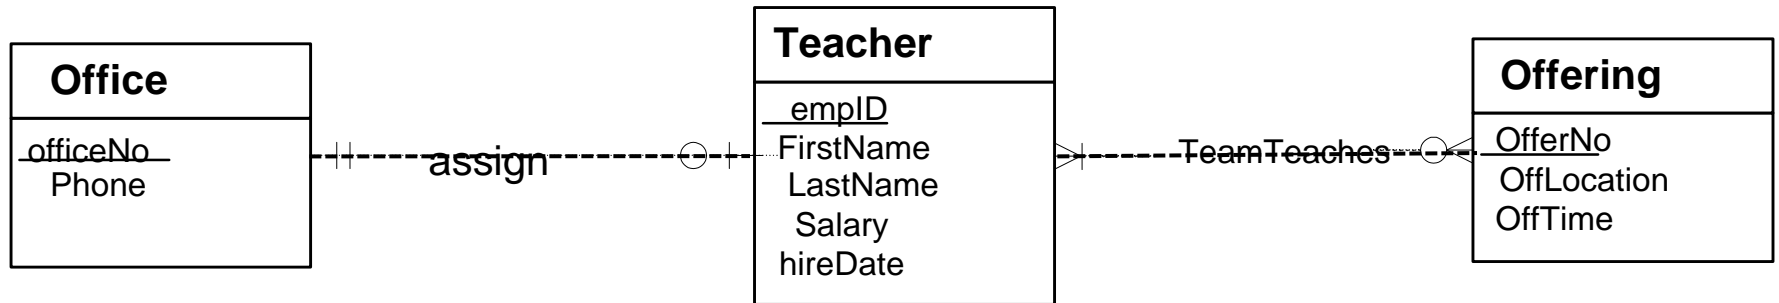
- Teacher:
 - Employee ID (ID), salary, firstname, lastname, hireDate
- Offering:
 - offerNo, Location, Time
- Multiple teachers can jointly teach one offering
- A teacher may or may not teach any offering, but a teacher can also teach many offerings
- One offering must be taught by at least one teacher

Example 3: Answer




- Multiple teachers can jointly teach one offering
- A teacher may or may not teach any offering, a teacher can also teach many offerings
- one offering must be taught by at least one teacher

Example 2 + Example 3



In MySQL Workbench ...

- An ***identifying relationship***: identified by a solid line between tables
 - An identifying relationship is one where the child table cannot be uniquely identified without its parent. Typically this occurs where an intermediary table is created to resolve a many-to-many relationship. In such cases, the primary key is usually a composite key made up of the primary keys from the two original tables.
- A ***non-identifying relationship***: identified by a broken (dashed) line between tables



In MySQL Workbench ...

- To edit relationship in MySQL Workbench,
 - Double click the relationship between tables, or
 - Right click the relationship between tables and select "Edit Relationship"

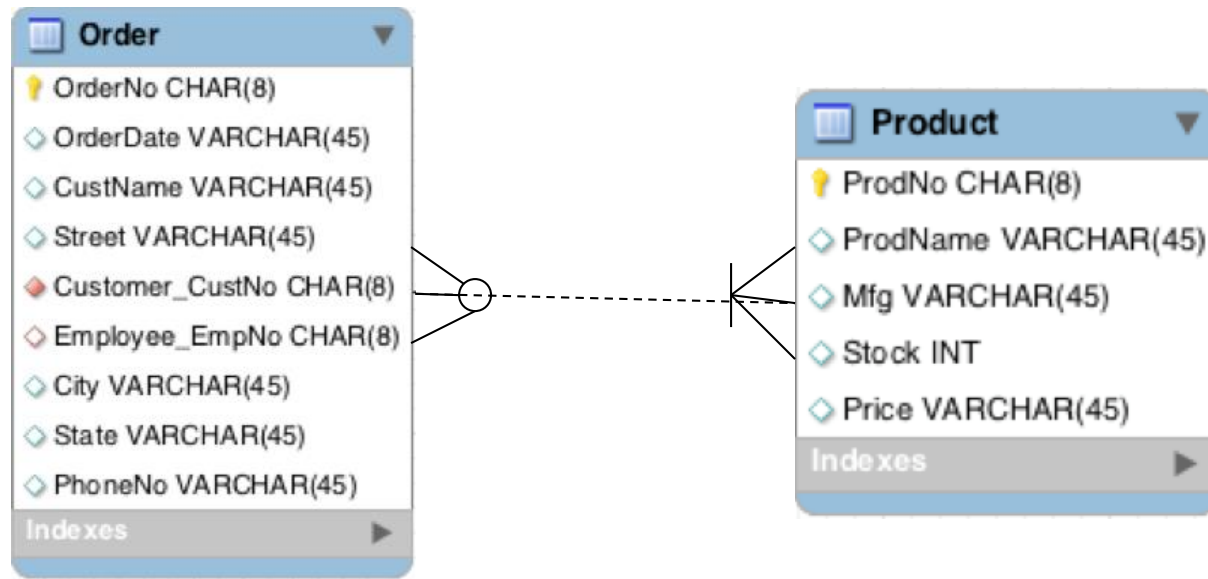
The screenshot shows the 'Edit Relationship' dialog box in MySQL Workbench. The dialog has a title bar 'Relationship' with a close button. It is divided into three main sections: 'Referencing Table', 'Cardinality', and 'Referenced Table'.
- The 'Referencing Table' section shows 'SALES_ORDER' with a foreign key 'fk_SALES_ORDER_ORDER' and a primary key 'ORDER_ORDER_NO: INT (PK)'. It has a 'Mandatory' checkbox checked and an 'Edit Table...' button.
- The 'Cardinality' section shows two options: 'One-to-One (1:1)' (selected with a radio button) and 'One-to-Many (1:n)' (unselected). There is an 'Invert Relationship' button and an 'Identifying Relationship' checkbox checked.
- The 'Referenced Table' section shows 'ORDER' with a primary key 'ORDER_NO: INT (PK)'. It has a 'Mandatory' checkbox checked and an 'Edit Table...' button.
At the bottom, there are two tabs: 'Relationship' and 'Foreign Key', with 'Foreign Key' currently selected.

Referencing Table	Cardinality	Referenced Table
SALES_ORDER Foreign Key: fk_SALES_ORDER_ORDER ORDER_ORDER_NO: INT (PK) <input checked="" type="checkbox"/> Mandatory Edit Table...	<input checked="" type="radio"/> One-to-One (1:1) <input type="radio"/> One-to-Many (1:n) Invert Relationship <input checked="" type="checkbox"/> Identifying Relationship	ORDER ORDER_NO: INT (PK) <input checked="" type="checkbox"/> Mandatory Edit Table...

Relationship Foreign Key

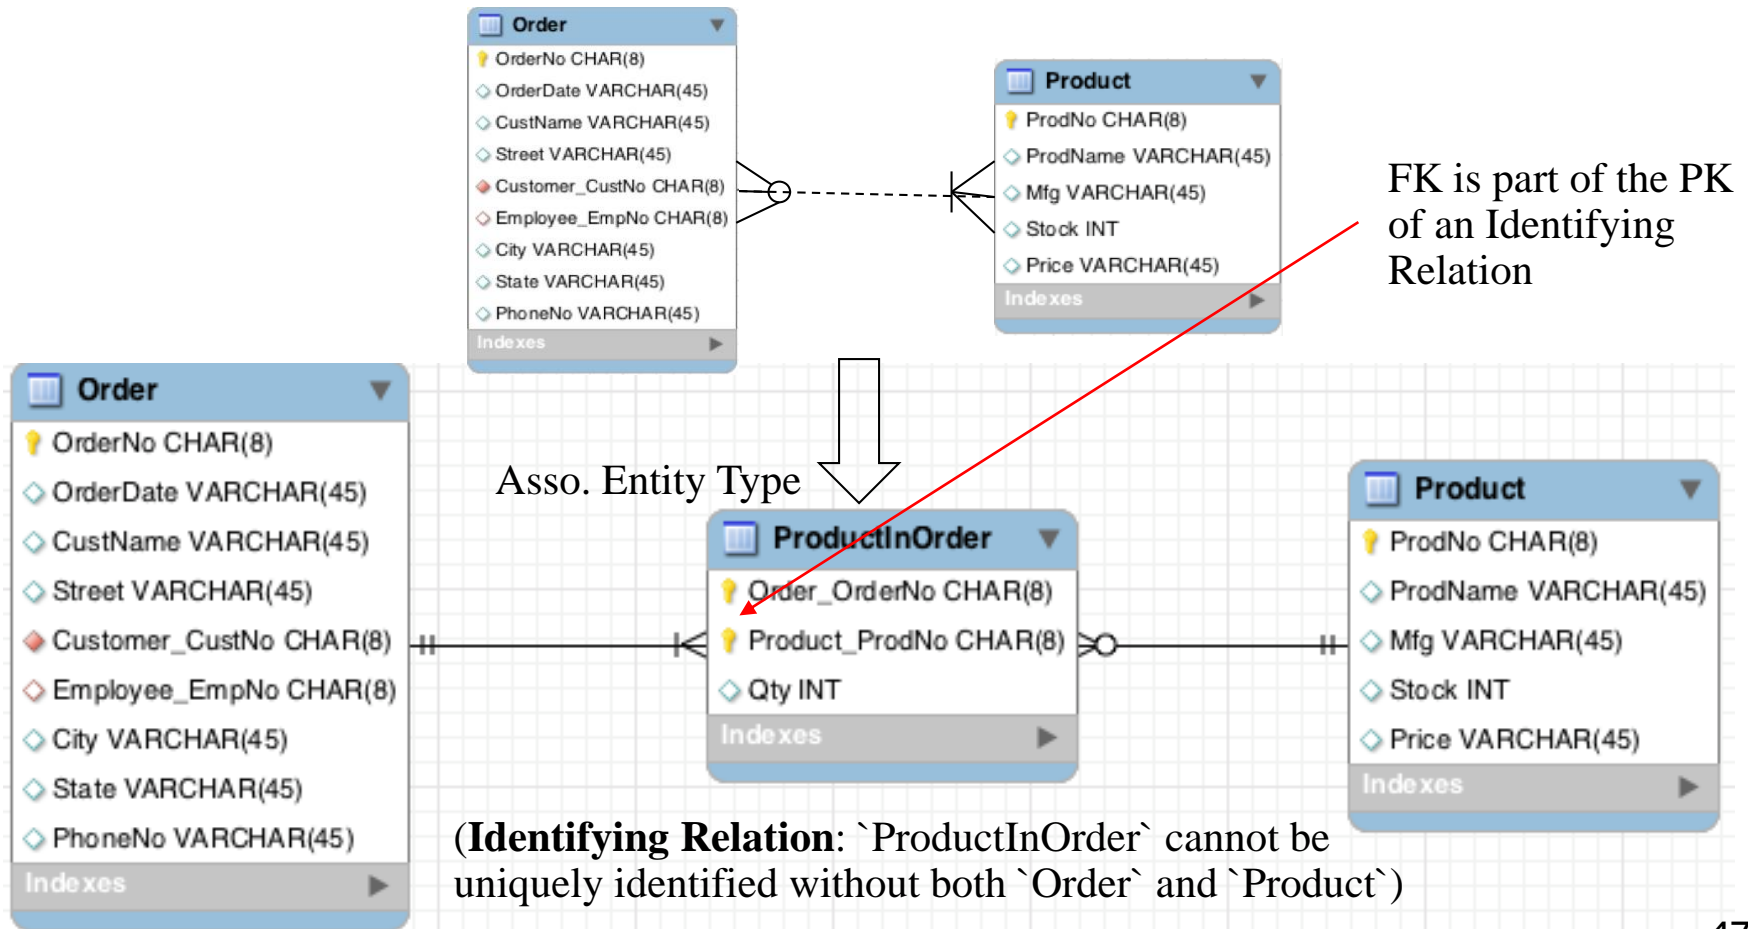
Implement M-N Relationships

- An *order* is associated with a number of *products* and a *product* can be associated with many *orders*.
- Each order should have at least one product. But a product may not have any order.
- The relationship between **Order** and **Product** is M:N



Implement M-N Relationships

- Transform M:N to 1:M Relationship:
 - M:N = two 1-M + associative entity type



Implement M-N Relationships

Tutorial Q2(a)

- A database involves the relationship between *parts*, the *vendors* providing the parts, and the *products* that use the parts. Specifically, a part is supplied by at least one vendor. Each vendor may supply many parts; but some new vendors may supply none. Each product is composed of one or more parts. A part may be used by many products; but some new parts may not be used by any product. Draw an entity-relationship diagram to illustrate the relationships among the entities **Part**, **Vendor**, and **Product**. In your solution, many-to-many relationships should be converted into one-to-many relationships. It is not necessary to provide attributes for the entities.